Augmenting LLM Code Translation with Compiler Analysis for C to Triton Kernel Generation

Xiao Qin
University of Leeds
Leeds, United Kingdom
ppnh4382@leeds.ac.uk

Chunwei Xia
University of Leeds
Leeds, United Kingdom

Zheng Wang
University of Leeds
Leeds, United Kingdom

ABSTRACT

Emerging programming models like Triton enable developers to better exploit modern accelerators, but translating legacy code to Triton remains challenging. While Large Language Models (LLMs) show promise for code translation, they often generate incorrect or suboptimal implementations due to a lack of precise parallelization reasoning. We present TritonPilot, a compiler-assisted LLM framework for translating C loop nests to Triton kernels. TritonPilot extracts parallelization properties using compiler-based analysis and encodes them as structured analysis facts in the LLM prompt, enabling correct and performance-aware parallel code synthesis. A profiling feedback loop further improves performance by feeding GPU hardware utilization metrics back to the LLM for iterative optimization. Across PolyBench/C and TSVC kernels, TritonPilot achieves 97% and 95% correctness respectively, improving over LLM-only baselines. At 8× problem sizes, the generated Triton kernels achieve a median 24.6× speedup over single-threaded C. Profiling feedback improves 9 of 29 PolyBench kernels at standard sizes and 12 of 21 at larger sizes, with gains up to 40× on individual kernels. We further demonstrate generalization on 8 Rodinia and ECP proxy application kernels (6/8 pass). These results show that combining compiler analysis with hardware profiling feedback enables both correct and high-performance LLM-based parallel code generation.

CCS CONCEPTS

• **Software and its engineering** → **Compilers**; • **Computing methodologies** → *Parallel computing methodologies*.

KEYWORDS

LLM code generation, GPU kernels, Triton, polyhedral analysis

ACM Reference Format:
Xiao Qin, Chunwei Xia, and Zheng Wang. 2026. Augmenting LLM Code Translation with Compiler Analysis for C to Triton Kernel Generation. In *Proceedings of Workshop on AI for High Performance Compiler Construction (AI4HPCC '26).* ACM, New York, NY, USA, 5 pages. https://doi.org/10.1145/nnnnnnn.nnnnnnn

AI4HPCC '26, July 6–9, 2026, Belfast, Northern Ireland, UK
2026. ACM ISBN 978-x-xxxx-xxxx-x/YY/MM. . . $15.00
https://doi.org/10.1145/nnnnnnn.nnnnnnn

1 INTRODUCTION

Accelerators such as GPUs are now essential for high-performance scientific computing, yet much existing software is still written in legacy C code that is not designed for parallel execution on modern accelerators. Translating these legacy programs to accelerator-friendly implementations is challenging, as it requires identifying parallelism, restructuring computation, and managing synchronization, tasks that require expertise in programming models and hardware.

Emerging programming models such as Triton [1] simplify GPU programming by providing higher-level abstractions for writing efficient kernels. However, mapping legacy C code to Triton still requires careful reasoning about loop-level parallelism, data dependencies, and execution order. Developers must determine which loop dimensions can be safely parallelized, how memory accesses should be organized, and when synchronization is required - decisions that are subtle and error-prone in practice.

Large Language Models (LLMs) have recently shown promise in code translation, including generating GPU kernels from high-level descriptions [2, 3]. However, LLMs lack precise reasoning about program semantics, especially data dependencies. As a result, they often produce incorrect or inefficient code, including unsafe parallelization that leads to data races, suboptimal kernel structures, and missing synchronization.

In this paper, we present TritonPilot, a system that augments LLM-based translation with compiler analysis. TritonPilot extracts parallelization properties from C loop nests using polyhedral analysis (PET) [4] and LLVM dependence analysis [5], and encodes them as structured guidance in the LLM prompt. This allows the model to generate correct and performance-aware Triton code without requiring it to perform dependence analysis itself.

We evaluate TritonPilot on PolyBench/C, TSVC, and application kernels on an NVIDIA RTX 3090 GPU. TritonPilot achieves 97% correctness on PolyBench/C and 95% on TSVC, improving over unguided baselines. At 8× problem sizes, the generated kernels achieve a median 24.6× speedup. We further demonstrate generalization on 8 Rodinia and ECP application kernels. This paper makes the following contributions:

- A compiler-assisted LLM framework for translating C loop nests to Triton, with a unified analysis module that derives all parallelization constraints from a single polyhedral analysis invocation and encodes them as structured facts.

- A profiling-guided optimization loop that feeds GPU hardware utilization metrics (from NVIDIA Nsight Compute) back to the LLM, enabling iterative performance improvement while preserving correctness.
- Empirical evaluation on 189 kernels across PolyBench/C, TSVC, Rodinia, and ECP benchmarks, showing that compiler analysis improves correctness and profiling feedback significantly improves performance.

2 OUR APPROACH

Figure 1 shows the TritonPilot workflow. Given a C source file, TritonPilot extracts loop nests, analyzes their parallelization properties, constructs a structured prompt, invokes an LLM to generate Triton code, and optionally refines performance using GPU profiling feedback. TritonPilot is implemented in approximately 14K lines of Python, using PET/isl [8] for polyhedral analysis and LLVM 17 [5] for fallback dependence analysis.

2.1 Unified Analysis

TritonPilot calls PET [4] once per kernel to extract a polyhedral representation (iteration domains, access relations, schedules as ISL objects). From this single invocation, it derives all parallelization constraints:

(i) Parallelization safety. For each loop dimension, TritonPilot checks for cross-iteration dependences by composing ISL access relations: given read map R and write map W to the same array, the conflict relation $R \circ W^{-1}$ is intersected with the iteration domain and the identity map is subtracted. A non-empty result marks the dimension as unsafe. *(ii) WAR dependences.* Write-after-read conflicts are identified when an array is both read and written at different iteration points, requiring array cloning or sequential execution. *(iii) Memory access patterns.* Stencil patterns (neighboring-element reads), triangular access (`j < i`), and cross-phase dependencies (multiple phases writing different arrays). *(iv) Reductions.* Scalar writes combined with array reads indicate reductions, mapped to Triton primitives (`tl.sum()`, `tl.dot()`).

When polyhedral analysis fails (e.g., non-affine accesses), TritonPilot falls back to LLVM's `DependenceAnalysis`, parsing direction vectors to determine which loop levels carry dependences.

2.2 Prompt Construction

All analysis results are stored in a unified JSON schema with fixed fields regardless of what patterns the kernel exhibits. A single, pattern-agnostic function renders the analysis into structured text. This design principle—*tell the LLM what is true about the kernel, rather than prescribing what to do*—improves scalability: adding a new analysis pass requires only extending the JSON schema, not modifying prompt templates. The prompt includes the C source, analysis facts (e.g., "`j` is safe to parallelize; `t` has cross-phase dependencies; stencil pattern detected"), the required function signature, and Triton coding rules.

2.3 Generation and Iterative Refinement

TritonPilot invokes the LLM (Claude Sonnet 4) and validates the generated kernel against the C reference using identical inputs and FP32 tolerance checks. On failure, errors are classified into four categories: compilation/runtime errors, numerical errors, missing synchronization barriers, and low performance ($<0.1\times$). Each triggers a targeted retry prompt. This repeats up to 10 attempts, with a context reset at attempt 6 to mitigate error accumulation.

2.4 Profiling-Guided Optimization

After correctness validation, TritonPilot optionally profiles the kernel with NVIDIA Nsight Compute (NCU), collecting SM throughput, memory throughput, occupancy, and L1/L2 cache hit rates. From these metrics, it classifies the bottleneck as *compute-bound* (high SM utilization), *memory-bound* (high memory throughput), or *latency-bound* (both low, indicating insufficient parallelism).

The bottleneck classification, raw metrics, and current code are fed back to the LLM with bottleneck-specific optimization guidance: memory-bound kernels are guided toward tiling and coalescing; latency-bound kernels toward increased grid size. The optimized kernel is re-validated for correctness and kept only if faster; otherwise, the previous version is retained. This loop runs up to 3 iterations, with NCU results cached when the implementation is unchanged.

3 EVALUATION SETUP

Platforms. All experiments are conducted on a GPU workstation equipped with an NVIDIA RTX 3090 GPU (24 GB memory, 936 GB/s memory bandwidth) and a dual-socket Intel Xeon Gold 5218R CPU (80 logical cores). We use PyTorch 2.5.1, Triton 3.1.0, CUDA 12.4, GCC 11.4.0, and Claude Sonnet 4 [7] as the underlying LLM.

Benchmarks. We evaluate on three benchmark suites: (1) *PolyBench/C* [6]: 30 numerical kernels (stencils, linear algebra, data mining) at default sizes (N=60–120) and 8× scale (N=480–960); (2) *TSVC* [15]: 151 loop kernels covering linear traversals, reductions, indirect addressing, and loop-carried dependences, benchmarked at 32 MB and 3200 MB per array; (3) 8 application kernels from Rodinia [13] (SRAD, lud, gaussian, lavaMD, hotspot, pathfinder) and ECP proxy applications [14] (miniMD LJ force, miniFE SpMV).

Baselines. The primary comparison is against *OMP GPU*—OpenMP GPU offloading compiled by NVIDIA `nvc` with managed memory, running on the *same* GPU as Triton. This isolates the quality of LLM-generated Triton code from hardware differences. We additionally report against *OMP CPU* (80-thread OpenMP) and *ST* (single-threaded GCC `-O3`).

Evaluation methodology. Speedups are measured over 50 iterations after 5 warmup iterations. Each iteration includes array re-initialization (GPU `clone()` for Triton, CPU copy for the C reference) to ensure in-place kernels start from identical data. Triton timing includes `cuda.synchronize()`.

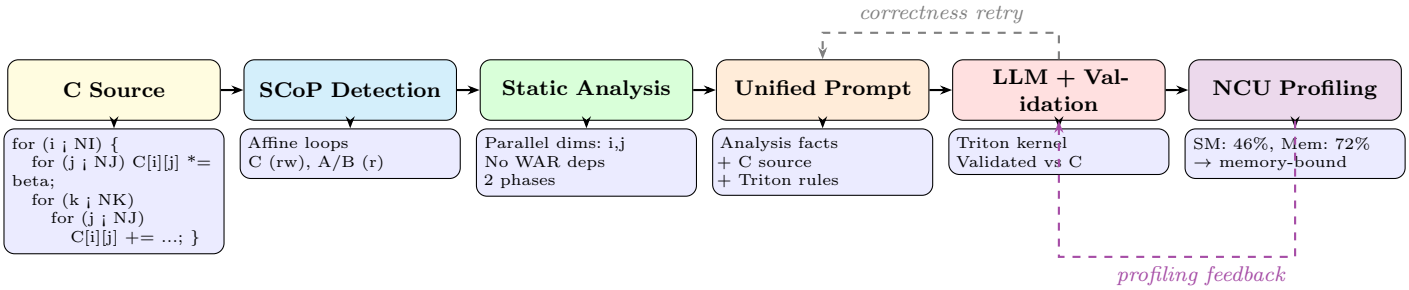

Figure 1: The TritonPilot workflow. Dashed arrows show iterative loops: the gray loop retries on correctness failures; the violet loop feeds NCU profiling metrics back to the LLM for performance optimization.

Table 1: PolyBench/C results. WA = with analysis, NA = no analysis.

Metric	WA	NA
Pass rate	29/30 (97%)	28/30 (93%)
First-try passes	12	8
Median speedup	0.76×	0.88×
Mean speedup	1.24×	1.64×
Kernels > 1×	11/29	12/28
Head-to-head (28 common):		
WA wins (>10% faster)	12	
NA wins (>10% faster)	9	
Ties (within 10%)	8	
WA-only pass	1 (jacobi_2d)	

For the correctness ablation, we compare *with analysis* (WA) vs. *no analysis* (NA, LLM receives only the C source). For profiling feedback, we report the initial speedup (first correct kernel) and the final speedup after up to 3 NCU-guided optimization iterations.

4 EXPERIMENTAL RESULTS

4.1 PolyBench/C (30 Kernels)

Table 1 summarizes results at standard (1×) problem sizes. With analysis guidance, 29/30 kernels pass correctness (1 failure due to LLM nondeterminism); without it, 28/30 pass but analysis uniquely enables `jacobi_2d`. The median speedup at 1× sizes is limited by PolyBench's small defaults ($N = 60$–120), where GPU launch overhead dominates. Compute-bound kernels show strong speedups (`3mm`: 4.01×, `covariance`: 3.66×, `deriche`: 3.58×, `2mm`: 2.97×).

In the head-to-head comparison, WA wins 12 of 28 common kernels while NA wins 9. The NA wins are largely due to LLM nondeterminism (different code quality across runs) rather than a systematic advantage of unguided generation.

Scaling. At 8× sizes (N: 60→480, matrices: 120→960), 27/30 pass with median speedup of 24.57× over single-threaded C, confirming that GPU advantage grows with problem size. Top kernels: `3mm` (2041×), `2mm` (2004×), `gemm` (451×), `syrk` (62×).

Table 2: Representative profiling feedback improvements (PolyBench). "Before" = first correct kernel; "After" = best after NCU-guided iterations.

Kernel	Before	After	Scale
1× sizes:			
3mm	0.10×	4.01×	1×
covariance	0.11×	3.66×	1×
2mm	1.48×	2.97×	1×
syrk	1.65×	2.07×	1×
8× sizes:			
gemm	120×	451×	8×
correlation	12×	51×	8×
atax	0.4×	15×	8×
syrk	24×	62×	8×

4.2 Profiling-Guided Optimization

Profiling feedback improved 9 of 29 kernels at 1× and 12 of 21 benchmarked kernels at 8× (Table 2). The gains are largest at 8× sizes where hardware utilization bottlenecks are more pronounced: `gemm` improves from 120× to 451× as the LLM restructures for better memory coalescing after receiving the memory-bound diagnosis. Correctness is preserved: each optimized kernel is re-validated, and the majority of optimization attempts that break correctness are automatically rejected.

4.3 TSVC (151 Kernels)

TritonPilot passes 143/151 TSVC kernels (94.7%), 91 on first attempt. At 32 KB–32 MB per array, LLM-generated Triton achieves a mean speedup of 2.78× over single-threaded C, with a maximum of 267×.

4.4 Application Kernels

We apply TritonPilot to 3 Rodinia and 5 ECP/Rodinia application kernels (Table 3). 6 of 8 pass correctness with analysis guidance. The highest speedups come from compute-intensive kernels: `LJ_force` (24.1×) and `SRAD` (12.8×). Two kernels fail: `hotspot` (complex 2D stencil with boundary conditions) and `lavaMD` (irregular particle interactions), both requiring domain-specific optimizations beyond what the unified analysis provides.

Table 3: Application kernels: Triton speedup vs. single-threaded C.

Source	Kernel	WA/ST	Pass?
ECP	LJ force	24.1×	Y
ECP	SpMV	7.1×	Y
Rodinia	SRAD	12.8×	Y
Rodinia	pathfinder	3.6×	Y
Rodinia	gaussian	0.6×	Y
Rodinia	lud	1.3×	Y
Rodinia	hotspot	—	N
Rodinia	lavaMD	—	N

5 RELATED WORK

LLM Code Generation for GPUs. Chen et al. [2] find frequent correctness issues in ChatGPT-generated CUDA kernels. KernelBench [3] reports 20–40% functional correctness for LLM-generated GPU kernels without guidance. Recent works improve generation through RL fine-tuning: Dr. Kernel [16] uses profiling-based rewards, Tehrani et al. [17] RL-fine-tune GPT-5, and DICE [18] applies RL to diffusion LLMs. Search-based methods keep the LLM frozen: AVO [19] uses LLM agents as evolutionary operators, K-Search [20] decouples planning from code instantiation, and OptiML [21] uses MCTS with profiler feedback. KernelBlaster [22] enables cross-task learning via persistent memory. For non-GPU targets, AscendKernelGen [23] and AscendCraft [24] address NPU kernels. These works focus on *how* to generate/optimize kernels; our work is orthogonal, addressing *what* parallelism is safe via compiler analysis. The two directions are complementary: compiler-derived constraints could serve as correctness rewards in RL frameworks or safety priors in search.

Polyhedral Compilation and Auto-Scheduling. PPCG [12] generates CUDA from affine loop nests but produces low-level code that lacks Triton's block-level optimizations. TVM [9] and Halide [10] use auto-scheduling with learned cost models. TritonPilot is complementary: it uses polyhedral analysis to extract *what* parallelism is safe, then delegates the *how* to the LLM, which can handle non-affine constructs, irregular access patterns, and Triton abstractions that polyhedral compilers cannot target directly.

6 CONCLUSION

We presented TritonPilot, a system that combines compiler-based analysis with LLM code generation for translating C loop nests to Triton GPU kernels. A unified analysis module calls PET once per kernel and derives all parallelization constraints (safe dimensions, WAR dependences, reductions, memory patterns) from the single polyhedral representation. These are rendered as structured facts in the LLM prompt via a pattern-agnostic renderer, replacing pattern-specific templates.

TritonPilot achieves 97% correctness on PolyBench/C (29/30) and 95% on TSVC (143/151), improving over unguided baselines (93% and lower on PolyBench). At 8× sizes, the generated Triton achieves a median 24.6× speedup. It generalizes to 6 of 8 Rodinia and ECP application kernels (up to 24×).

A profiling feedback loop using NCU hardware metrics improved 9 of 29 PolyBench kernels at 1× and 12 of 21 at 8×, with gains up to 40×. However, the LLM frequently breaks correctness during optimization, confirming that single-shot prompting is insufficient. Future work will integrate compiler analysis as a correctness oracle and NCU metrics as a performance reward within search-based or RL-based optimization frameworks.

REFERENCES

[1] P. Tillet, H. T. Kung, and D. Cox. Triton: An intermediate language and compiler for tiled neural network computations. In *Proc. MAPL*, 2019.

[2] Y. Chen, S. Huang, D. Zheng, et al. ChatGPT for CUDA programming: Evaluating LLM-generated GPU code. *arXiv preprint arXiv:2404.12000*, 2024.

[3] A. Ouyang, S. Zheng, et al. KernelBench: Can LLMs write efficient GPU kernels? *arXiv preprint arXiv:2404.12489*, 2024.

[4] S. Verdoolaege and T. Grosser. Polyhedral extraction tool. In *Proc. IMPACT*, 2012.

[5] C. Lattner and V. Adve. LLVM: A compilation framework for lifelong program analysis & transformation. In *Proc. CGO*, 2004.

[6] L.-N. Pouchet. Polybench/C: The polyhedral benchmark suite, version 4.2.1. https://web.cs.ucla.edu/~pouchet/software/polybench/, 2016.

[7] Anthropic. Claude Sonnet 4 model card. https://www.anthropic.com/claude, 2025.

[8] S. Verdoolaege. isl: An integer set library for the polyhedral model. In *Proc. ICMS*, 2010.

[9] T. Chen, T. Moreau, Z. Jiang, et al. TVM: An automated end-to-end optimizing compiler for deep learning. In *Proc. OSDI*, 2018.

[10] J. Ragan-Kelley, C. Barnes, A. Adams, et al. Halide: A language and compiler for optimizing parallelism, locality, and recomputation in image processing pipelines. In *Proc. PLDI*, 2013.

[11] U. Bondhugula, A. Hartono, J. Ramanujam, and P. Sadayappan. A practical automatic polyhedral parallelizer and locality optimizer. In *Proc. PLDI*, 2008.

[12] S. Verdoolaege, J. Carlos Juega, A. Cohen, et al. Polyhedral parallel code generation for CUDA. *ACM TACO*, 9(4), 2013.

[13] S. Che, M. Boyer, J. Meng, et al. Rodinia: A benchmark suite for heterogeneous computing. In *Proc. IISWC*, 2009.

[14] M. A. Heroux, D. W. Doerfler, P. S. Crozier, et al. Improving performance via mini-applications. Sandia National Laboratories Technical Report SAND2009-5574, 2009.

[15] D. Callahan, J. Dongarra, and D. Levine. Vectorizing compilers: A test suite and results. In *Proc. Supercomputing*, 1988.

[16] W. Liu, J. Xu, Y. Li, et al. Dr. Kernel: Reinforcement learning done right for Triton kernel generations. *arXiv preprint arXiv:2602.05885*, 2026.

[17] A. Tehrani, Y. Emara, E. Wissam, et al. Fine-tuning GPT-5 for GPU kernel generation. *arXiv preprint arXiv:2602.11000*, 2026.

[18] H. Bai, L. Kong, X. Chen, et al. DICE: Diffusion large language models excel at generating CUDA kernels. *arXiv preprint arXiv:2602.11715*, 2026.

[19] T. Chen, Z. Ye, B. Xu, et al. AVO: Agentic variation operators for autonomous evolutionary search. *arXiv preprint arXiv:2603.24517*, 2026.

[20] S. Cao, Z. Mao, J. E. Gonzalez, and I. Stoica. K-Search: LLM kernel generation via co-evolving intrinsic world model. *arXiv preprint arXiv:2602.19128*, 2026.

[21] A. Bhattacharjee, H. Ping, S. V. Le, et al. OptiML: An end-to-end framework for program synthesis and CUDA kernel optimization. *arXiv preprint arXiv:2602.12305*, 2026.

[22] K. S. Dong, S. Modi, D. Nikiforov, et al. KernelBlaster: Continual cross-task CUDA optimization via memory-augmented in-context reinforcement learning. *arXiv preprint arXiv:2602.14293*, 2026.

[23] X. Cao, J. Zhai, P. Li, et al. AscendKernelGen: A systematic study of LLM-based kernel generation for neural processing units. *arXiv preprint arXiv:2601.07160*, 2026.

[24] Z. Wen, S. Shao, Z. Li, et al. AscendCraft: Automatic Ascend NPU kernel generation via DSL-guided transcompilation. *arXiv preprint arXiv:2601.22760*, 2026.